INTRODUCTION AND OPTIMIZATION

CELERY SCHEDULE STRATEGY

INTRODUCTION

Celery is a simple, flexible, and reliable distributed system to process vast amounts of messages, while providing operations with the tools required to maintain such a system.

It's a task queue with focus on real-time processing, while also supporting task scheduling.

docs.celeryproject.org

INTRODUCTION

- Simple
  Celery is easy to use and maintain, and it *doesn't need configuration files*.

- Highly Available
  Workers and clients will automatically retry in the event of connection loss or failure, and some brokers support HA in way of *Primary/Primary* or *Primary/Replica* replication.

- Fast
  A single Celery process can process millions of tasks a minute, with sub-millisecond round-trip latency (using RabbitMQ, librabbitmq, and optimized settings).

- Flexible
  Almost every part of *Celery* can be extended or used on its own, Custom pool implementations, serializers, compression schemes, logging, schedulers, consumers, producers, broker transports, and much more.

CELERY WORKFLOW

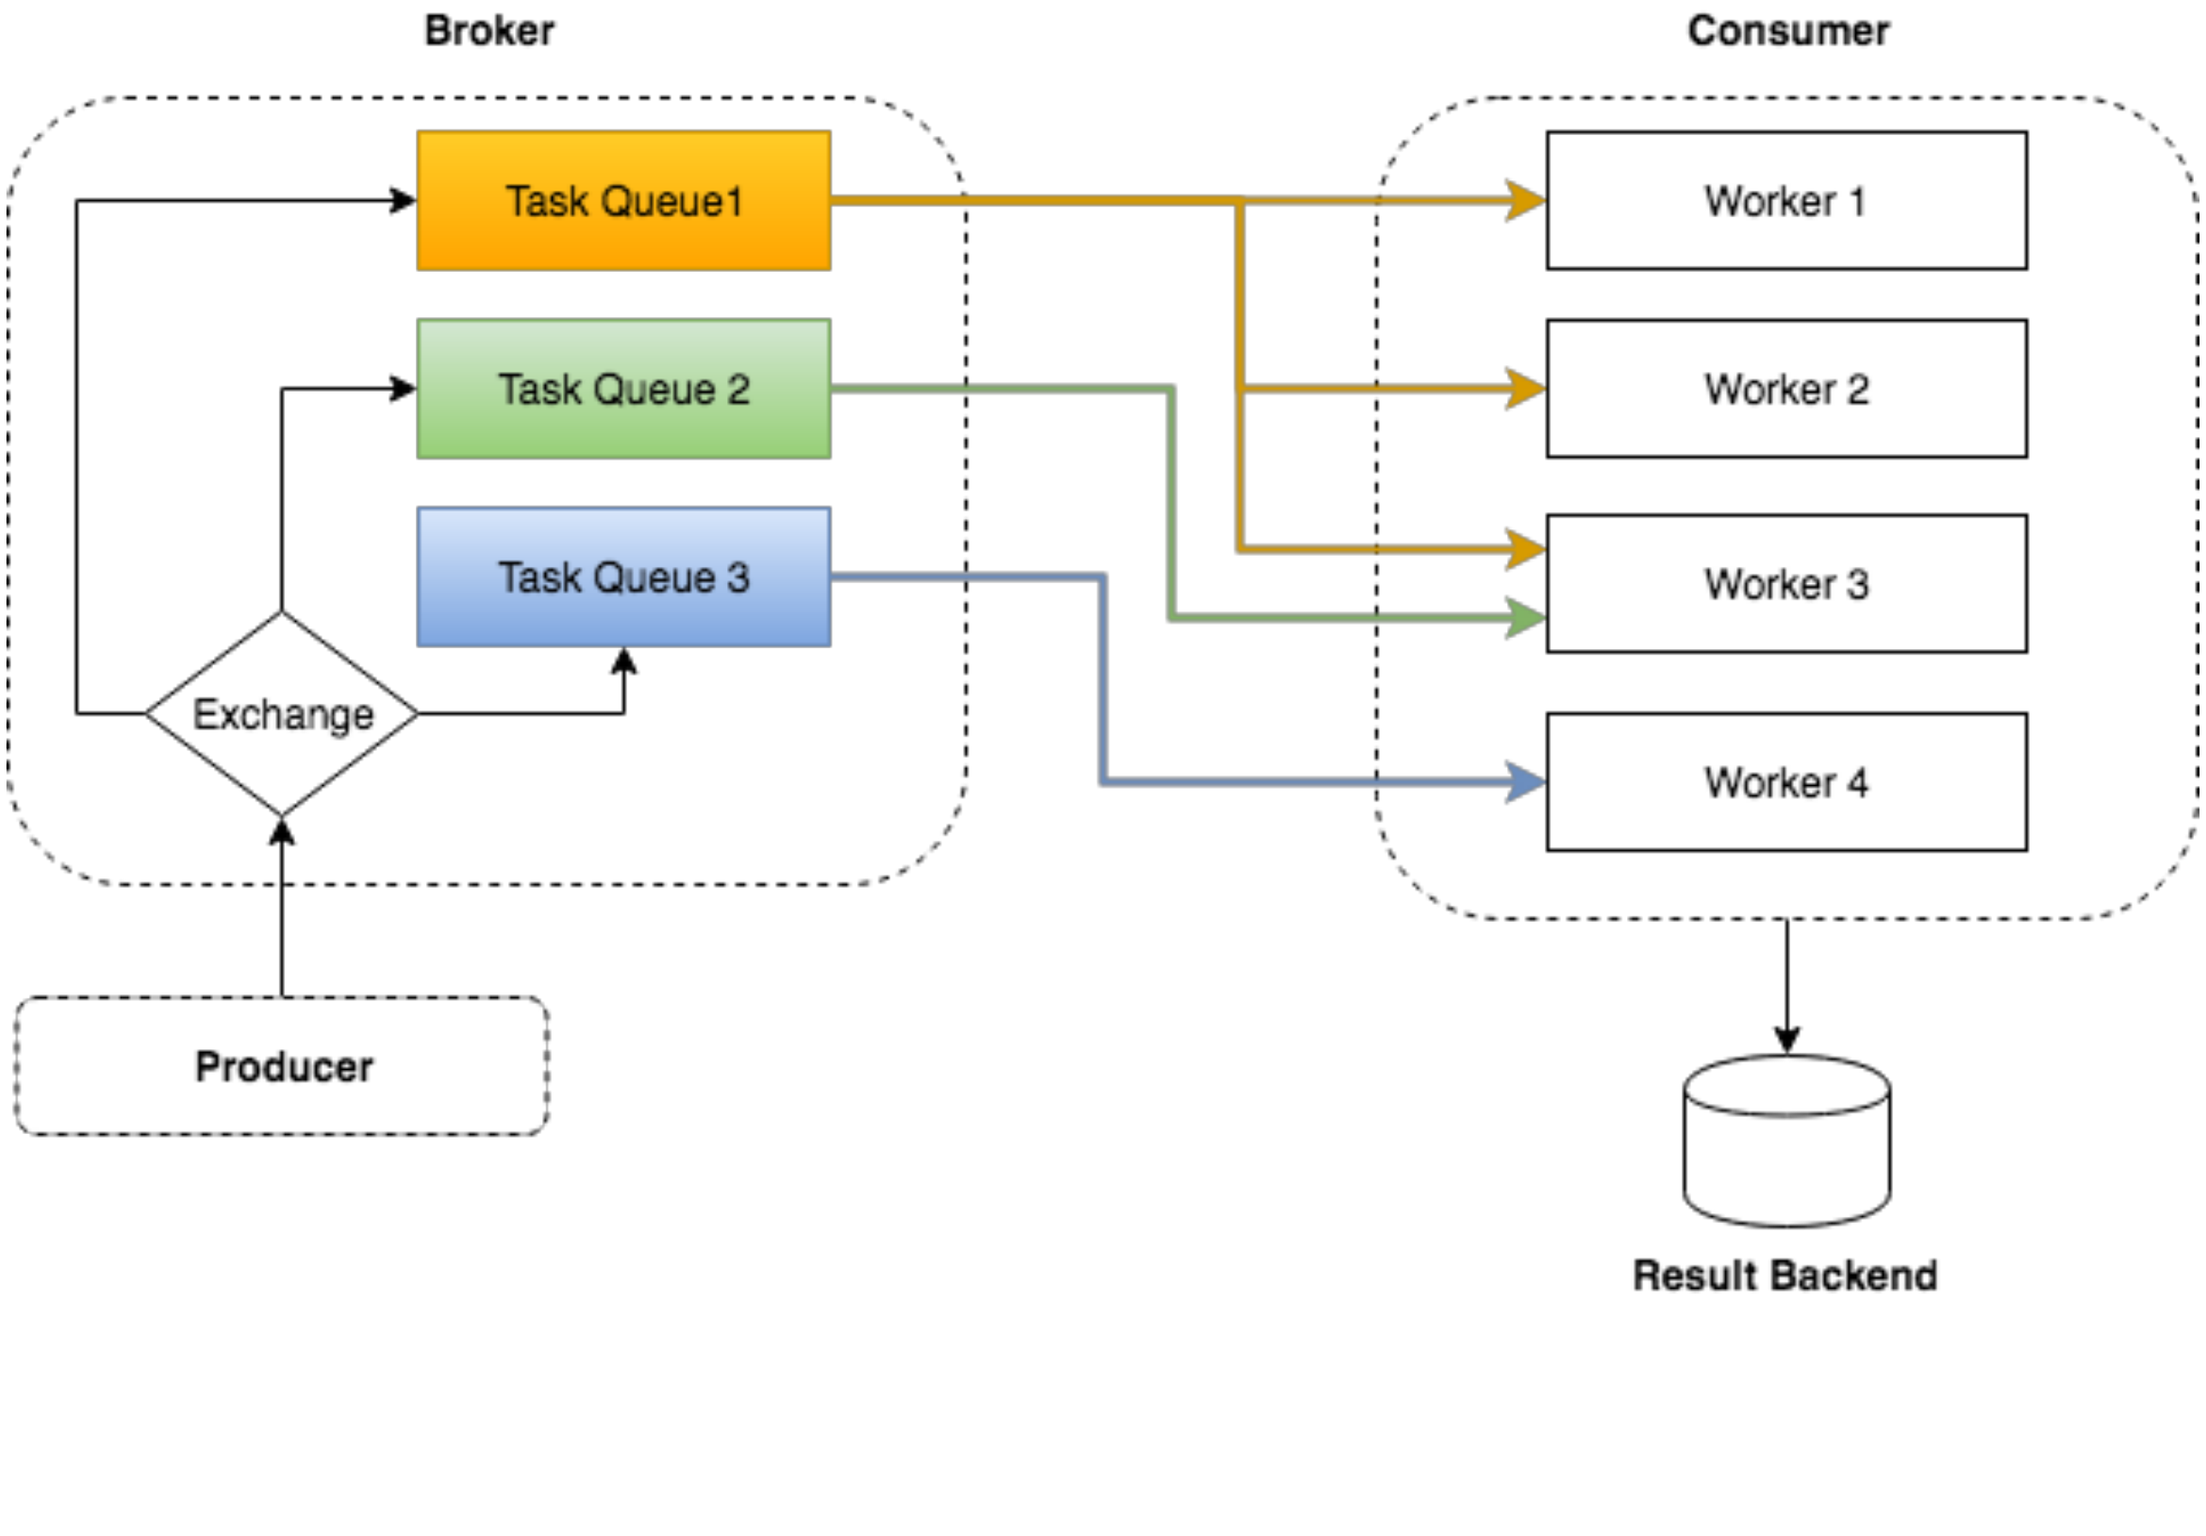

PREFETCH MULTIPLIER

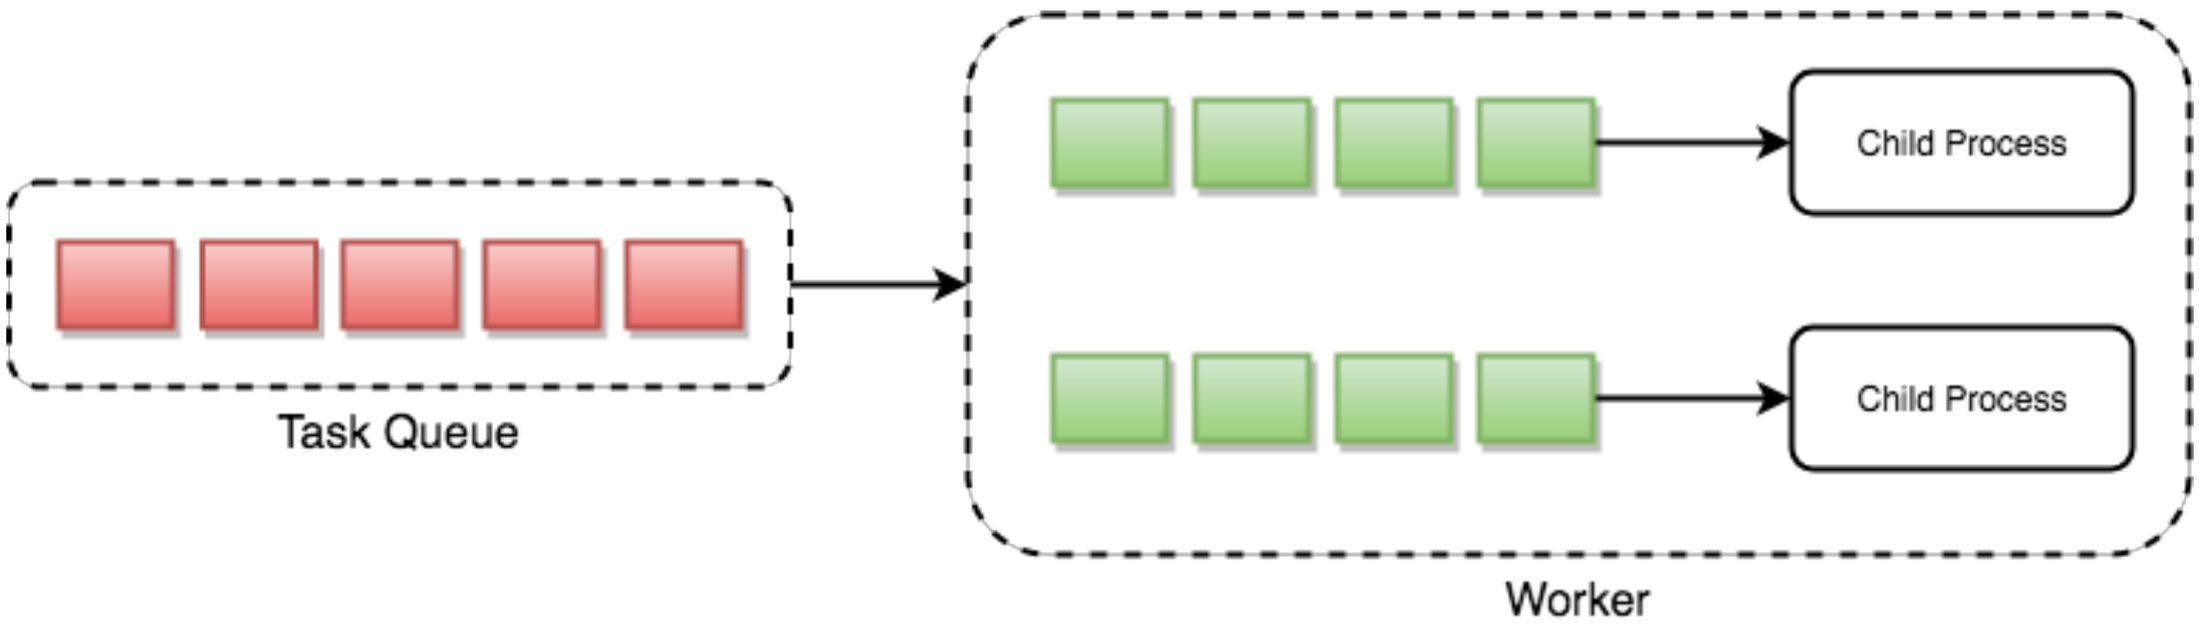

worker_prefetch_multiplier = 4

PREFETCH MULTIPLIER

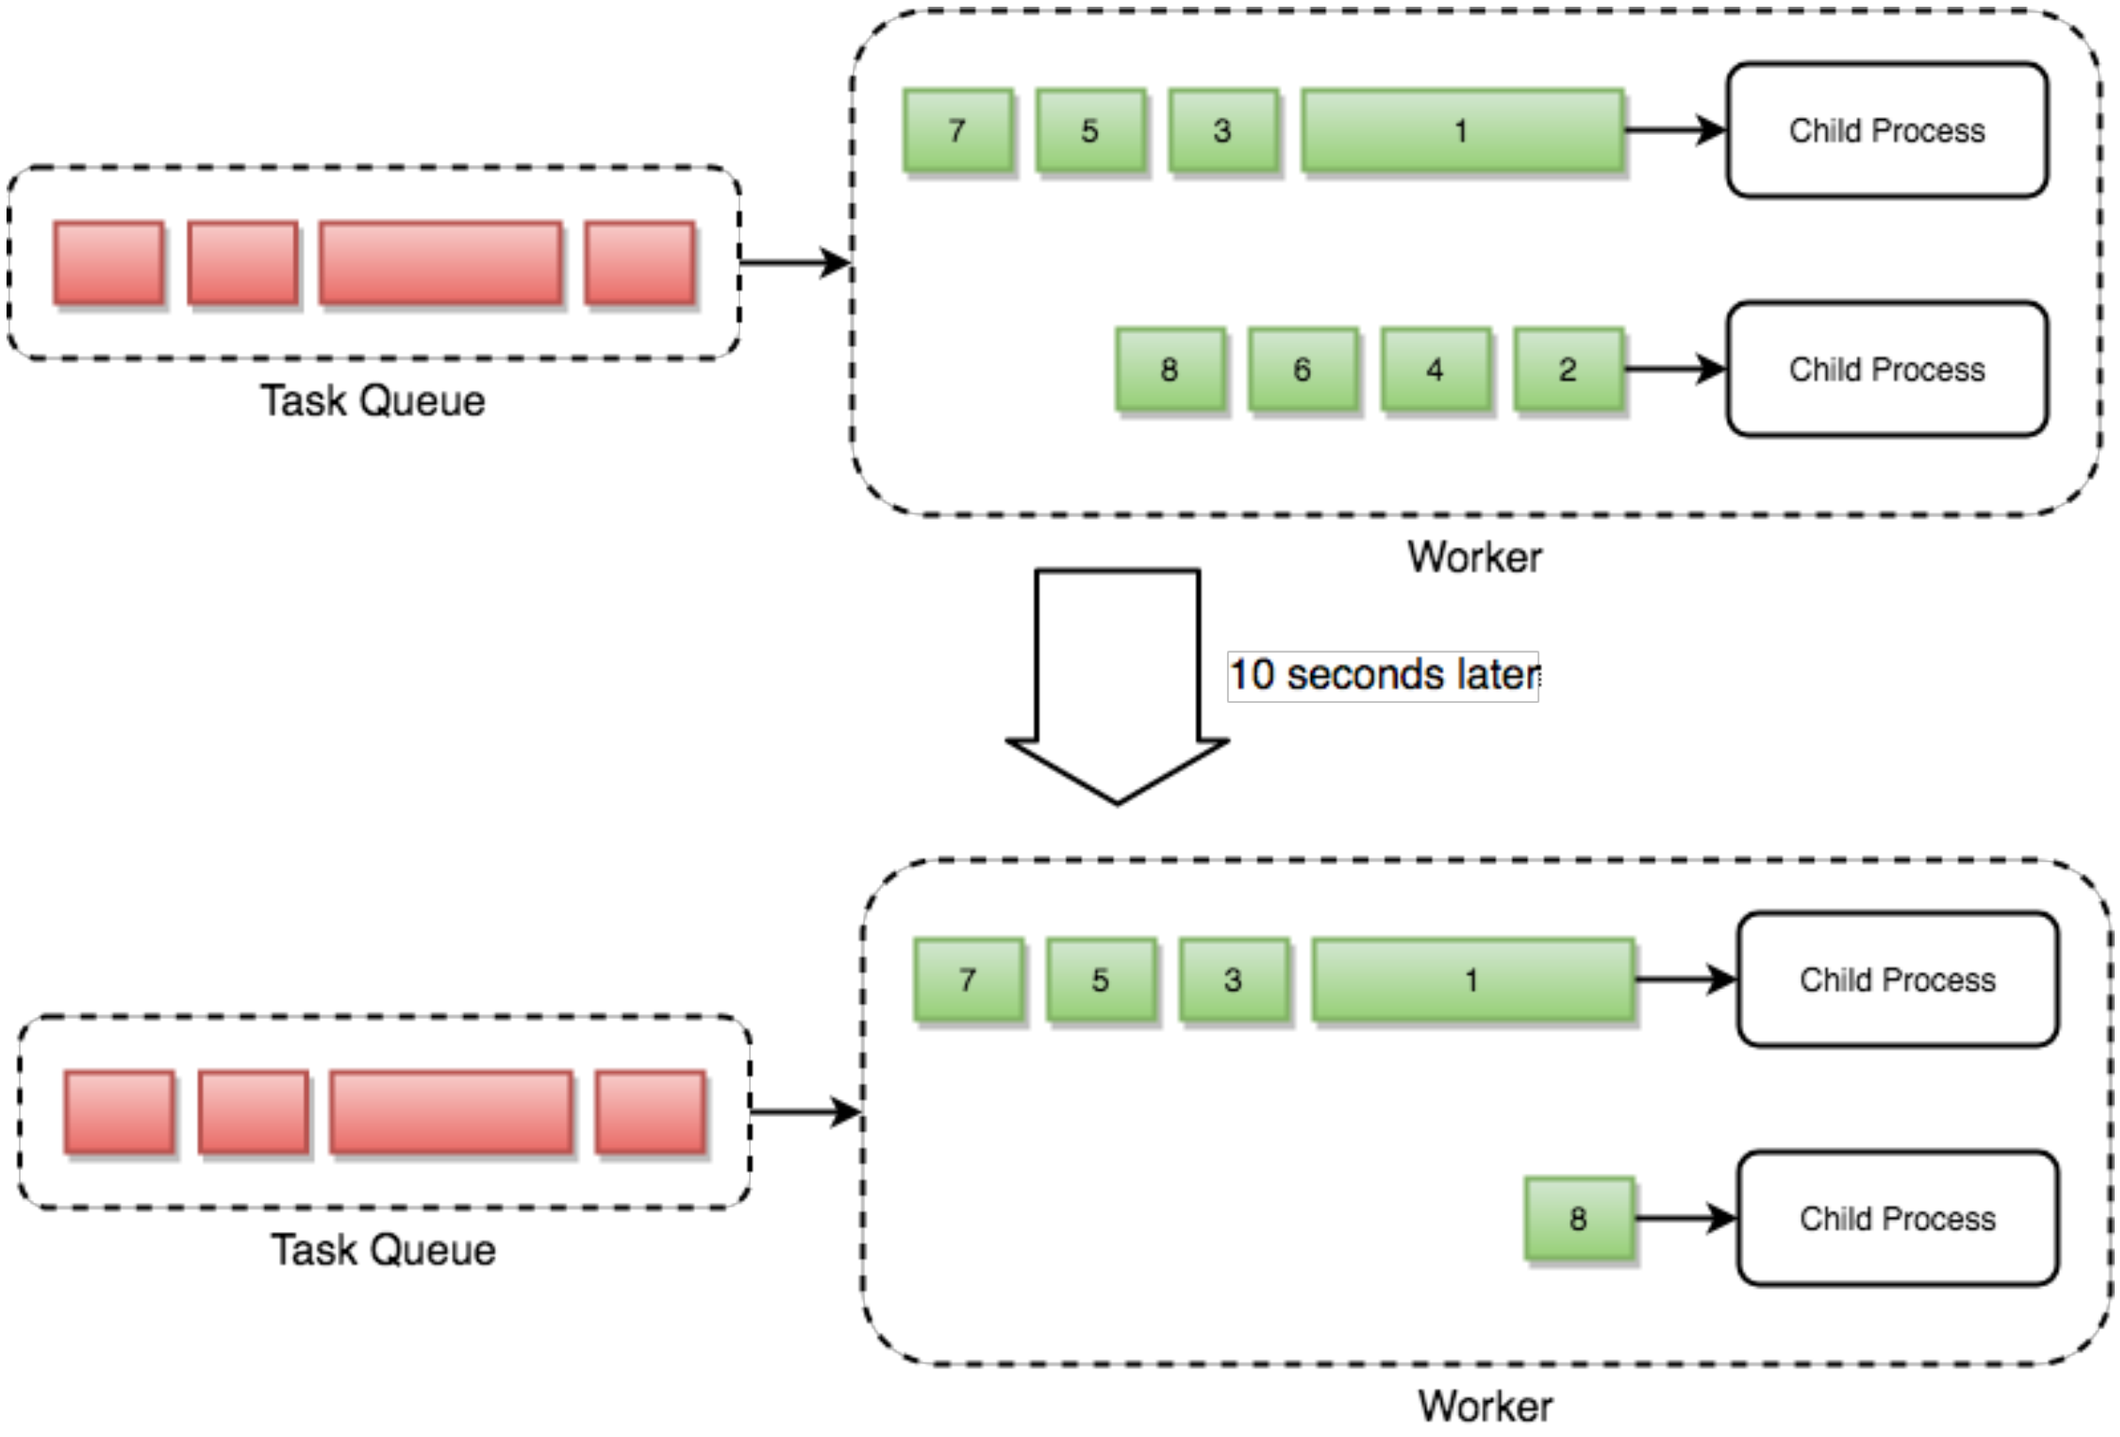

PREFETCH LIMIT

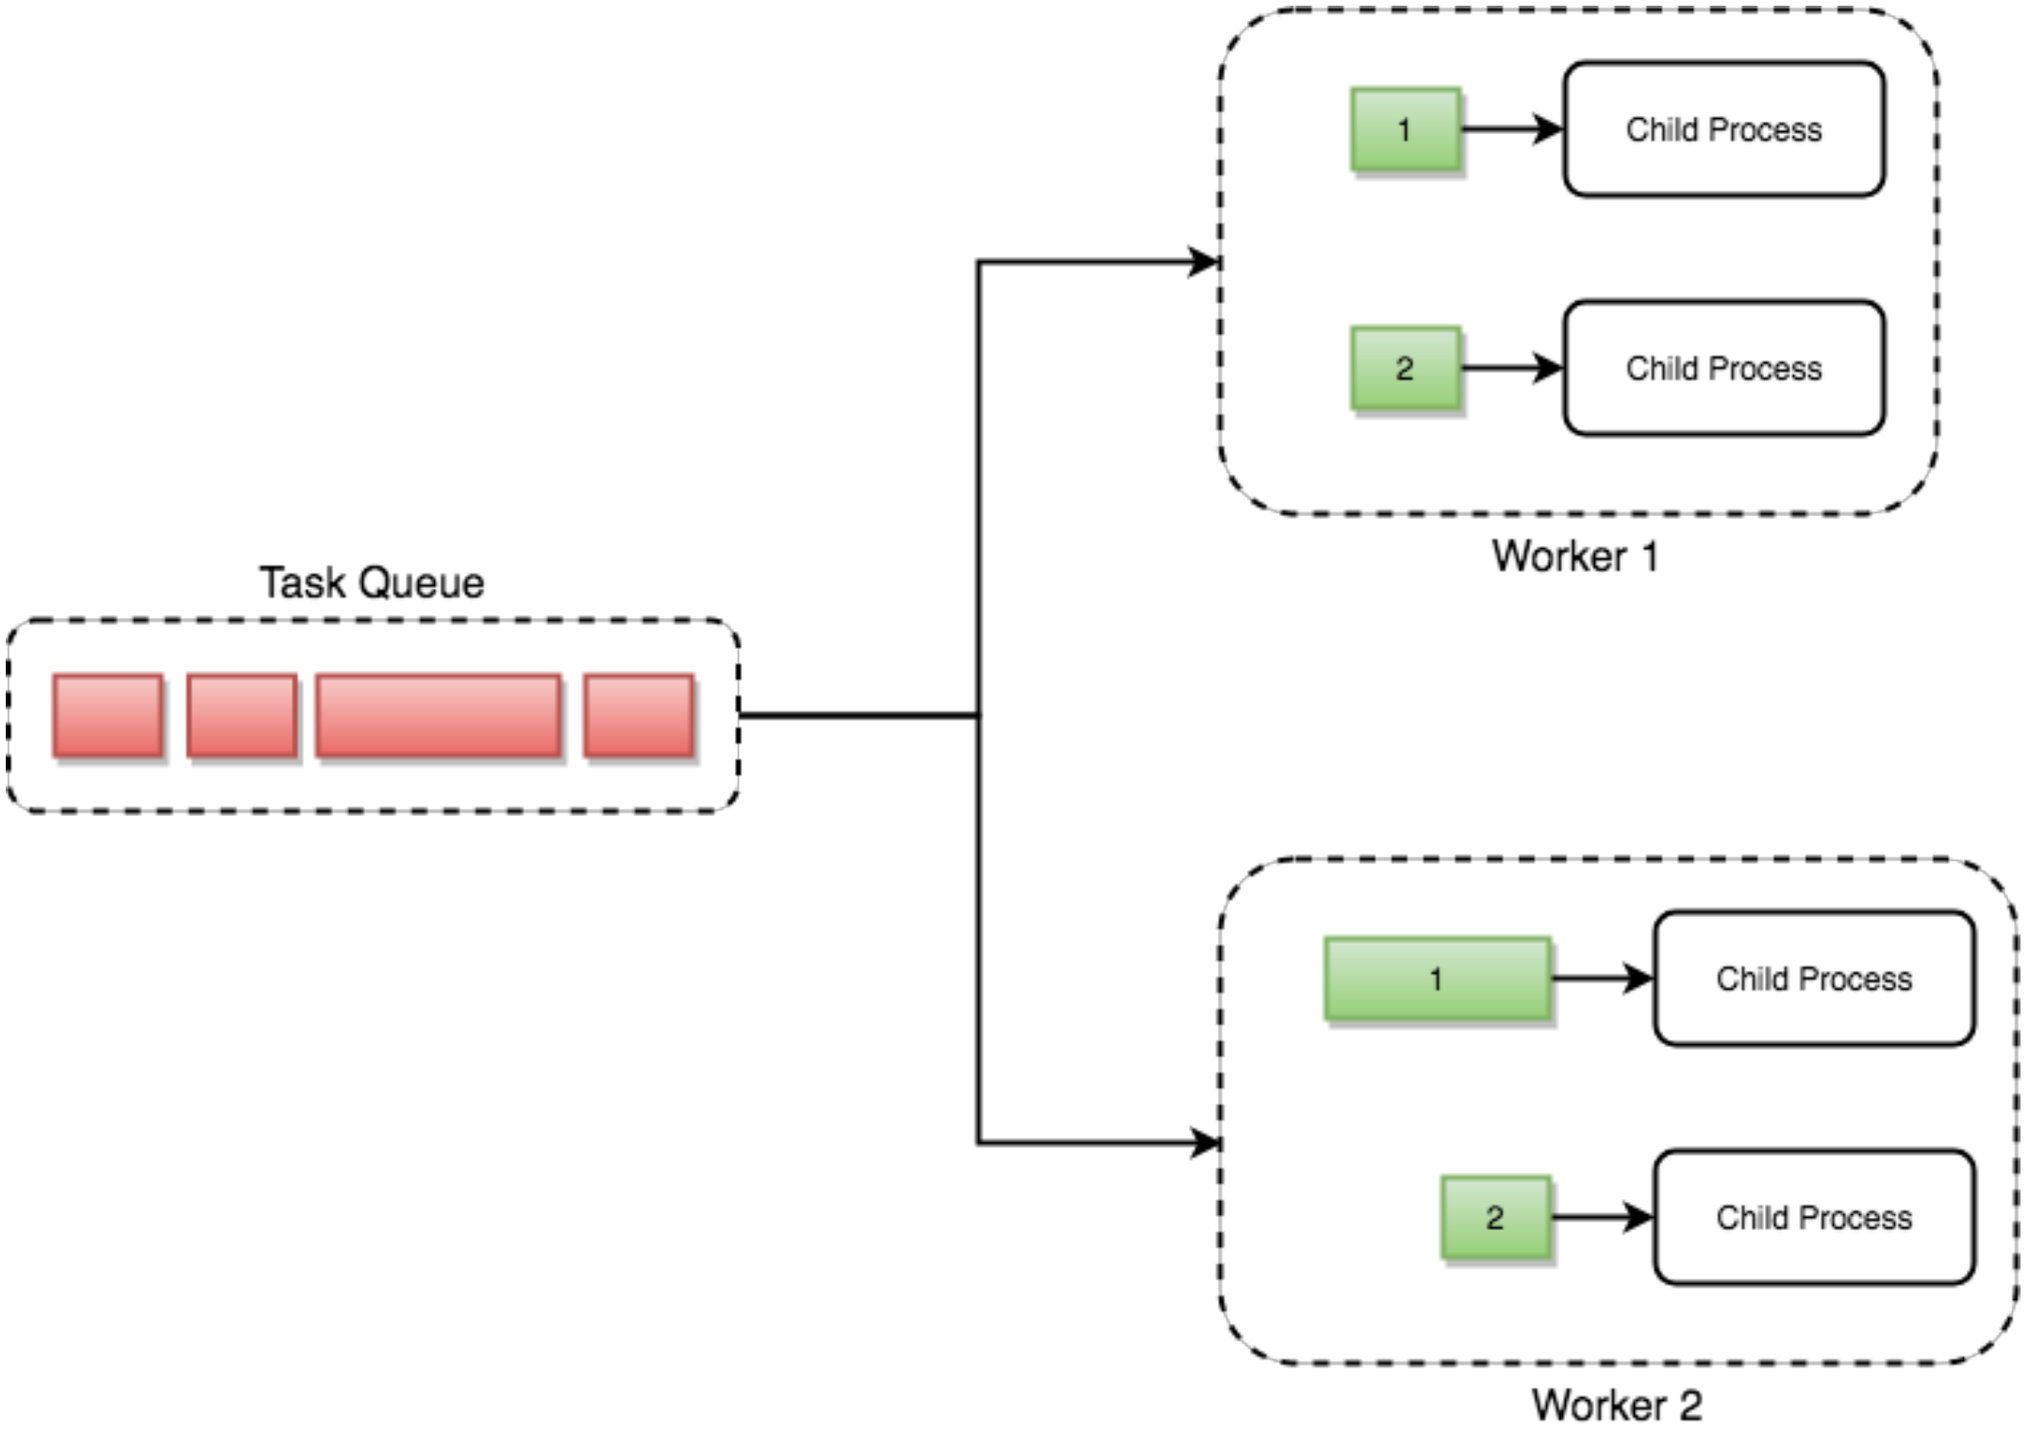

FAST VS FAIR

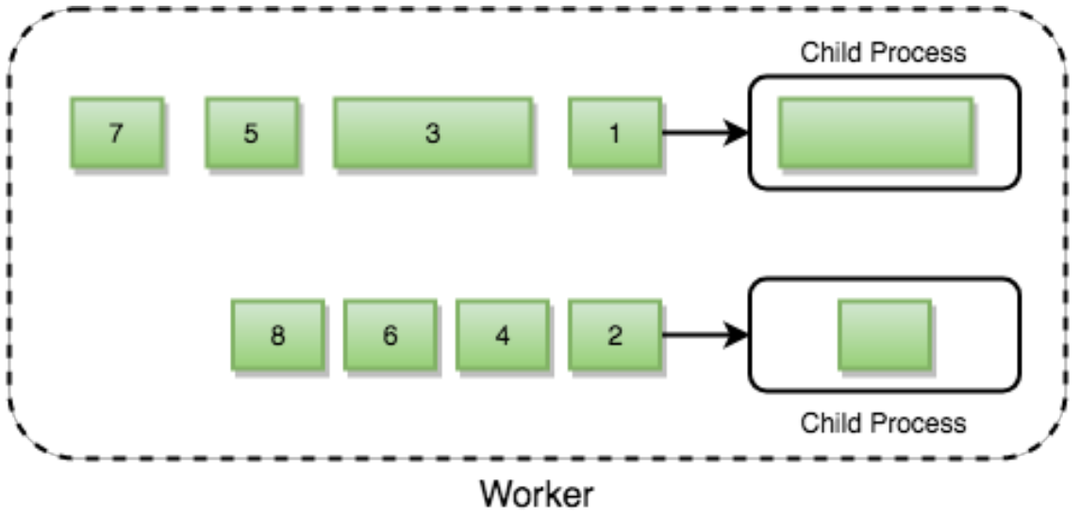

"fast" schedule strategy

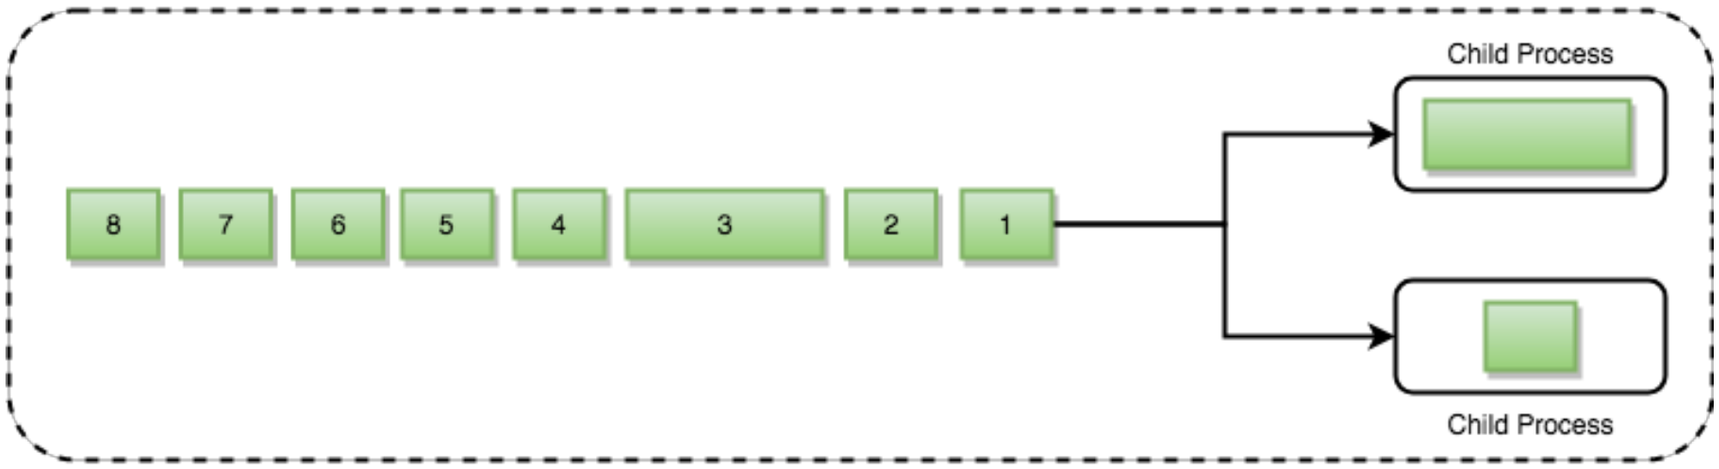

"fair" schedule strategy